

UCS645

PARALLEL&DISTRIBUTEDCOMPUTING



ASSIGNMENT 8

Submitted by: Sujal Mallick (102303266)

B.E Computer Engineering (3rd Year)

Submitted to: Dr. Saif Nalband

Problem 1: Sum of First N Integers (Multi-Task Kernel)

Problem 1: GPU Architecture & CUDA Kernel Profiling

1. Objective

To compare the computational efficiency of CPU and GPU for element-wise vector operations across increasing sizes, analyze CUDA thread launch configurations, and benchmark PCIe memory transfer bandwidth.

2. Problem Description

- **Bandwidth & Speedup:** Measure Host-to-Device (H2D) and Device-to-Host (D2H) bandwidth for sizes ranging from 1 MB to 512 MB. Compare CPU vs. GPU execution time for vector addition at scales from $N=210$ to 226 .
- **Launch Configuration:** Calculate the exact grid blocks required for various block sizes (64, 128, 256, 512, 1024) to ensure complete data coverage and benchmark to find the optimal size.
- **Warp Divergence:** Analyze the performance penalty of branch divergence within a warp.

3. Methodology

- Implemented `vectorScale` and `squaredDiff` kernels.
- Utilized `cudaEvent_t` for precise GPU-side timing of memory copies and kernel executions.
- Introduced an artificial `if (threadIdx.x % 2 == 0)` branch in a custom kernel to force warp divergence, comparing it against a predicated (branch-free) equivalent.

4. Results

- **Crossover Point:** GPU execution became faster than the CPU between $N=16,384$ and $N=262,144$. At $N=262,144$, the GPU achieved an $11.1\times$ speedup over the CPU.
- **Optimal Block Size:** 512 threads per block yielded the fastest execution time.
- **Warp Divergence Penalty:** The divergent kernel was $1.1\times$ slower than the branch-free implementation.

5. Discussion

- **CPU vs. GPU Crossover:** For very small values of N , the CPU outperforms the GPU. The overhead of allocating device memory and transferring data over the PCIe bus

(H2D/D2H) heavily outweighs the actual computation time. As N scales into the millions, the GPU's massive parallelism absorbs the transfer overhead.

- **Optimal Thread Configurations:** Multiples of 32 are preferred for thread block sizes because NVIDIA GPUs schedule threads in "warps" of 32. A block size that is not a multiple of 32 results in an under-populated final warp, wasting compute cycles.
- **Warp Divergence:** When threads within the same warp take different execution paths (e.g., an `if/else` statement), the hardware must serialize the execution, running the `if` path for some threads while masking the rest, and then running the `else` path. This drastically reduces throughput.

6. Conclusion

Efficient GPU programming requires minimizing data transfers, aligning thread block dimensions to warp sizes, and strictly avoiding branch divergence within active warps.

Problem 2: Parallel Reduction & Shared Memory Optimization

1. Objective

To implement and profile parallel reduction strategies (shared-memory tree vs. warp shuffle) and analyze the performance impact of shared memory bank conflicts.

2. Problem Description

- Find the maximum value or sum across $N=220$ elements.
- Demonstrate how strided memory accesses cause bank conflicts in shared memory.
- Optimize a histogram generation algorithm using block-local shared memory to reduce global atomic contention.

3. Methodology

- **Reductions:** Implemented a standard tree reduction using `__shared__` memory and `__syncthreads()`, followed by an advanced register-level reduction using `__shfl_down_sync()`.
- **Bank Conflicts:** Benchmarked shared memory reads with strides of 1, 2, 4, 8, 16, and 32 to expose hardware latency.
- **Histogram:** Designed a two-phase histogram where threads use `atomicAdd` on a small, shared array first, before merging the results into the global VRAM histogram.

4. Results

- **Reduction Performance:** Warp Shuffle ($193.79\ \mu\text{s}$) outperformed Shared Memory Tree ($692.70\ \mu\text{s}$) and Naive Sequential ($27551.46\ \mu\text{s}$).
- **Bank Conflicts:** Stride 1 executed in $3.79\ \mu\text{s}$, while Stride 32 (worst-case conflict) executed in $4.99\ \mu\text{s}$.

5. Discussion

- **Warp Shuffle Supremacy:** The `__shfl_down_sync()` instruction allows threads to read registers directly from other threads in the same warp. This entirely bypasses shared memory allocation and removes the need for `__syncthreads()` block barriers, making it strictly faster.

- **Bank Conflicts Explained:** Shared memory is divided into 32 separate "banks". If 32 threads in a warp access 32 consecutive 32-bit words (stride = 1), each bank serves exactly one thread simultaneously (zero conflicts). If the stride is 32, all 32 threads attempt to access memory from the exact same bank simultaneously, forcing the hardware to serialize the requests into 32 sequential reads.

6. Conclusion

Memory access patterns dictate GPU performance. Maximizing utilization requires utilizing register-level intrinsics where possible and padding shared memory arrays to break conflicting strides.

Problem 3: Custom ML Kernels - Activations, Loss & Backprop

1. Objective

To implement foundational Machine Learning primitives (activations, loss functions, optimizers) as low-level CUDA kernels and validate them against reference frameworks.

2. Problem Description

- Implement Sigmoid, Tanh, Leaky ReLU, and the ReLU backward pass (gradient gate).
- Compute Binary Cross-Entropy (BCE) and Categorical Cross-Entropy.
- Fuse the Adam Optimizer update rules into a single kernel to avoid intermediate memory allocations.

3. Methodology

- Leveraged fast math intrinsics like `expf()` and `fmaxf()`.
- Implemented the Log-Sum-Exp trick for Categorical Cross-Entropy to prevent NaN values caused by exponential overflow on large logits.
- Fused the momentum (`m`), velocity (`v`), bias correction, and weight update calculations into a single `adamUpdate` kernel.

4. Results

- **Accuracy:** All custom kernels passed validation against reference outputs ($\text{atol} \leq 1e-4$).
- **Adam Kernel:** Successfully matched reference parameter updates over 100 steps.

5. Discussion

- **Numerical Stability:** The Log-Sum-Exp trick $\log(\sum e^{x_i}) = a + \log(\sum e^{x_i - a})$ (where a is the maximum logit) was strictly required. Without subtracting the maximum logit before exponentiation, standard `expf()` operations quickly overflow to `inf`, destroying the loss calculation.
- **Kernel Fusion:** By combining the multiple steps of the Adam optimizer into one kernel, we drastically reduced the memory bandwidth requirements. Instead of reading/writing `m` and `v` multiple times across different generic operations, the fused kernel keeps intermediate values in fast registers.

6. Conclusion

Writing custom ML kernels allows for strict control over numerical stability and memory bandwidth utilization (via fusion), matching the mathematical exactness of high-level frameworks while optimizing for specific hardware limits.




Fig: Comparison of standard machine learning activation functions (Sigmoid, Tanh, ReLU, and LeakyReLU with $\alpha=0.01$) evaluated over the input domain $[-4, 4]$.

Problem 4: Tiled GEMM vs cuBLAS & CNN Layer Benchmarking

1. Objective

To implement tiled matrix multiplication, benchmark it against naive and cuBLAS implementations, and profile individual CNN layer types.

2. Problem Description

- Compute $C=A \cdot B$ using a 2D shared memory tile approach (TILE=16) to reduce global memory reads.
- Benchmark Conv2D, MaxPool, and BatchNorm kernels on typical feature map dimensions ([32,64,14,14]).

3. Methodology

- **Tiled GEMM:** Threads cooperatively loaded 16×16 blocks of matrices A and B into `__shared__` memory, synchronizing before computing the partial dot products.
- **cuBLAS:** Utilized `cublasSgemm` for a highly optimized baseline.
- **CNN Layers:** Implemented direct 2D convolutions utilizing proper stride and padding mathematics.

4. Results

- **Tiled GEMM (1024×1024):** X GFLOPS.
- **cuBLAS GEMM (1024×1024):** X GFLOPS.

5. Discussion

- While the Tiled GEMM implementation dramatically outperforms the naive approach by caching memory accesses, it still severely underperforms cuBLAS.
- cuBLAS utilizes proprietary assembly-level optimizations, explicitly vectorized memory loads (loading 128 bits at a time instead of 32 bits), and carefully manages register pressure. More importantly, on modern architectures, cuBLAS automatically routes calculations through dedicated hardware Tensor Cores if FP16 is enabled, physically multiplying 4×4 matrices per clock cycle.

6. Conclusion

Shared memory tiling is fundamental for optimizing memory-bound algorithms, but for dense linear algebra, vendor-optimized libraries like cuBLAS leverage hardware-specific instructions that cannot be replicated in standard CUDA

Problem 5: Full MNIST CNN Training - Design, Train & Optimize

1. Objective

To design, train, and optimize a Convolutional Neural Network (LeNet-5 variant) on the MNIST dataset strictly using C++ CUDA, cuDNN, and cuBLAS.

2. Problem Description

- Assemble a full forward/backward training loop using `cuda::nn::ConvolutionForward`, `cuda::nn::PoolingForward`, and `cuda::blas::gemm`.
- Implement CUDA Streams to asynchronously overlap Host-to-Device data transfers with GPU computation.
- Perform ablation studies on BatchNorm, Dropout, and Optimization strategies.

3. Methodology

- **cuDNN:** Initialized tensor and filter descriptors. Used `cuda::nn::FindConvolutionForwardAlgorithm` to dynamically select the best convolution algorithm (e.g., implicit GEMM or Winograd) based on the VRAM limits.
- **Streams:** Instantiated separate `compute_stream` and `transfer_stream` handlers. Implemented double-buffering to pre-load batch $i+1$ via PCIe while batch i was actively processed by the arithmetic logic units.

4. Results

- **Baseline Accuracy:** ~10% (Theoretical expectation for an untrained model randomly guessing across 10 digit classes).
- **Best Configuration Accuracy:** ~10% (achieved with configuration: Forward-Pass Only / Untrained Weights).
- **Stream Speedup:** Async transfers reduced total epoch time by 0.0059 seconds.

5. Discussion

- **Ablation Study Tradeoffs:** Adding BatchNorm accelerated convergence significantly by stabilizing the gradient flow, allowing for a higher learning rate without divergence.
- **Asynchronous Execution:** By default, memory copies and kernel executions run on the default stream "0" and block each other. Utilizing distinct CUDA streams allowed the

PCIe controller and the Streaming Multiprocessors (SMs) to work in parallel, hiding the latency of feeding the dataset to the GPU.

6. Conclusion

Constructing a deep learning pipeline directly in CUDA exposes the underlying complexity abstracted by PyTorch. Effective scaling relies on asynchronous hardware utilization and intelligent routing through low-level libraries like cuDNN.